




RESEARCHMIND

AI Research Intelligence Platform

BUSINESS FOUNDATION DOCUMENT

MPA 2026 — Project Deliverable

Teodor Mihailescu	Adela Danescu	Anastasia Sandu	Selena Hurloi
Lead Dev & AI Arch.	Backend & QA	Business & Product	Design & Pitch

Table of Contents

1. MOTIVATION.....	3
2. SUMMARY	3
User Stories	4
3. PROPOSED SOLUTION DETAILS.....	4
3.1 THE THREE-AGENT ARCHITECTURE	4
3.2 SWOT ANALYSIS	5
3.3 MARKET ANALYSIS	5
3.4 TECHNOLOGY STACK	7
3.5 RISK REGISTER	8
4. PLANNING	9
4.1 TEAM MEMBERS.....	9
4.2 RACI MATRIX	9
4.3 WORK PACKAGES.....	10
WP1 — Foundation & Architecture (Sprint 1, Weeks 1–2) Owner: Teodor Mihailescu.....	10
WP2 — AI Agent Development (Sprint 2, Weeks 3–4) Owner: Teodor Mihailescu	10
WP3 — Frontend & Integration (Sprint 3, Weeks 5–6) Owner: Teodor / Selena.....	11
WP4 — Business Docs & Pitch (Sprint 4, Weeks 7–8) Owner: Anastasia / Selena.....	12
4.4 SPRINT PLAN	13
4.5 GANTT CHART.....	14
5. COSTS	15
5.1 COST CATEGORIES & ESTIMATES	15
5.1.1 TOKEN COST ECONOMICS — GROQ INFERENCE PRICING	16
5.2 REVENUE MODEL & PROJECTIONS.....	17
5.3 COST/BENEFIT ANALYSIS (ROI & PAYBACK PERIOD)	17
5.4 GO-TO-MARKET & MARKETING STRATEGY	18
5.4.1 TARGET PERSONAS	18
5.4.2 CHANNEL STRATEGY & LAUNCH PLAN.....	19
5.4.3 MARKETING BUDGET & KPIs.....	20
6. ONE-PAGE SUMMARY — LEAN CANVAS	21
7. AI USAGE DISCLOSURE	21

1. Motivation

Academic research is fundamentally broken in one specific way: knowledge is siloed in individual papers, and the cognitive labour of connecting, synthesising, and generating new ideas from a growing literature rests entirely on the human researcher. The average PhD student reads 200–400 papers per year, takes notes in disconnected applications, and struggles to maintain a coherent mental model of their field. Studies show researchers spend up to 6 hours per week just searching for relevant literature — time that could be spent on actual discovery.

ResearchMind was built to solve this exact problem. We believe the next breakthrough in any field will come not from reading faster, but from connecting ideas smarter. Our platform uses a three-agent AI architecture to transform a collection of static papers into a living, searchable knowledge graph — one that gets smarter as you add more papers, surfaces hidden connections automatically, and proposes novel research directions no single researcher could see alone. A third agent critically evaluates each paper's methodology, statistical soundness, and reproducibility — helping researchers immediately identify which studies to trust and which to approach with caution.

The timing is right: large language models are now capable of understanding scientific text with near-expert accuracy, and free inference APIs (Groq) make real-time AI processing accessible without a large compute budget. ResearchMind is our answer to the question: "What if your entire research library could think — and think critically?"

2. Summary

ResearchMind is a SaaS platform for researchers that uses three specialised AI agents to transform academic paper libraries into interactive knowledge graphs. Researchers upload papers (or paste text), and ResearchMind automatically extracts concepts, methodology, findings, and keywords — then connects papers sharing similar ideas into a visual graph. A second AI agent synthesises the full library, surfaces research gaps, and generates novel, evidence-backed hypotheses. A third agent critically evaluates each paper's quality, flagging methodological weaknesses, biases, and reproducibility issues.

The product targets a three-tier market: individual researchers and PhD students (free and €29/month tiers), academic labs (team plan at €99/month), and universities seeking institutional licences (€499/month). Revenue scales naturally with the growth of the academic community, creating a strong network effect.

User Stories

ID	As a...	I want to...	So that...
US-01	Researcher	add a paper by pasting its URL or text	Agent 1 extracts structured knowledge automatically
US-02	Researcher	view a knowledge graph of my papers	I can see concept connections at a glance
US-03	Researcher	run synthesis across my full library	Agent 2 surfaces themes, gaps, and novel hypotheses
US-04	Researcher	critique any individual paper	Agent 3 scores its quality and flags weaknesses
US-05	Researcher	chat with my paper library	I get grounded answers without re-reading every paper
US-06	Researcher	see critique results saved per paper	I can compare quality across my collection

3. Proposed Solution Details

3.1 The Three-Agent Architecture

ResearchMind orchestrates three independent AI agents, each with a distinct responsibility:

- Agent 1 — Paper Analyzer: Ingests raw paper text (up to 5,000 characters). Extracts title, summary, key concepts, methodology, datasets used, main findings, limitations, and keywords as structured JSON. Temperature: 0.15 (deterministic).
- Agent 2 — Synthesis Agent: Reads all Agent 1 outputs together. Identifies specific cross-paper connections, research gaps, and generates 3 novel testable hypotheses citing paper titles by name. Temperature: 0.62 (creative).
- Agent 3 — Paper Critic: Evaluates a single paper's methodology rigour, statistical soundness, reproducibility, novelty, and clarity. Returns a credibility score (0–10) with score breakdown, strengths, weaknesses, and suggested improvements. Temperature: 0.35 (balanced).

3.2 SWOT Analysis

STRENGTHS	WEAKNESSES
◆ Three-agent AI — unique in the market ◆ Groq free tier: 700 tok/s, no credit card needed ◆ Paper Critic is unmatched by any competitor ◆ Open-source LLaMA 3.1 — zero per-token cost at MVP ◆ Academic network effect grows graph automatically	⚠ No brand recognition vs. Semantic Scholar ⚠ PDF parsing deferred to v1.1 ⚠ In-memory store — no persistence across restarts ⚠ Groq free tier: 14,400 requests/day limit
OPPORTUNITIES	THREATS
→ 5M+ academic papers published annually → EU Horizon 2027 mandates open-science tooling → AI research assistant market: \$2.8B by 2028 → University licence deals: 1 deal = thousands of users	- X Elsevier / Springer could build similar product - X OpenAI / Anthropic research-focused releases - X GDPR/IP concerns around storing paper text - X Groq could change free-tier pricing

3.3 Market Analysis

The academic research software market reached \$8.2B globally in 2025, growing at 12% CAGR (Grand View Research). The AI-powered sub-segment is the fastest-growing vertical, projected to reach \$2.8B by 2028, fuelled by the explosion of LLM capabilities post-2023. The landscape has shifted dramatically in 2024–2025 with the entry of Google NotebookLM, Perplexity Research Mode, and SciSpace — yet none of these tools combines multi-agent analysis, cross-paper synthesis with testable hypothesis generation, and per-paper quality critique in a single free pipeline.

Key differentiators for ResearchMind: (1) Three-agent pipeline — Analyze → Synthesize → Critique — with no equivalent in the market; (2) Paper Critic generates a 0–10 credibility score with dimension breakdown, unique in the space; (3) Groq free tier eliminates marginal inference cost up to ~21,000 full workflows/month; (4) Real-time D3.js knowledge graph with force-directed concept clustering.

Feature	Research Mind	SciSpace	Elicit	NotebookLM	Perplexity Research	Scholarcy	Semantic Scholar
3-Agent AI Pipeline	✓	X	X	X	X	X	X
Hypothesis Generator	✓	X	X	X	X	X	X
Paper Critic / Score	✓	X	X	X	X	X	X
Cross-paper Synthesis	✓	X	✓	✓	~	X	X
Knowledge Graph	✓ Live	✓	X	X	X	X	✓ Citations
Chat with Library	✓	✓	~	✓	✓	X	X
PDF / URL Ingestion	✓ URL	✓	✓	✓	✓	✓	✓
Free Tier	✓ Full	✓ Ltd	✓ Ltd	✓ Ltd	✓ Ltd	~	✓ Search
Offline / Self-hosted	~	X	X	X	X	X	X
API Access	✓ Lab	✓	X	X	X	X	✓ Ltd
Open Source Stack	✓	X	X	X	X	X	X

Legend: ✓ = full support, ~ = partial/limited, X = not available. SciSpace and Elicit added major 2024–2025 updates but remain synthesis-only. NotebookLM excels at Q&A from uploaded PDFs but has no hypothesis engine, no critic, and no graph. Perplexity Research Mode is search-native, not library-native. ResearchMind is the only free tool with all four pillars: analysis + synthesis + critique + graph.

Differentiator: ResearchMind is the only free tool that combines paper analysis, cross-library synthesis with hypothesis generation, and quality critique with a numerical score in a single integrated pipeline — with real-time knowledge graph visualisation.

3.4 Technology Stack

Layer	Technology	Rationale
AI Inference	Groq API + LLaMA 3.1 8B Instant	700 tok/s; free up to 14,400 req/day; no credit card
Agent 1	Paper Analyzer (Python/Groq)	Extracts title, concepts, methodology, findings, keywords
Agent 2	Synthesis Agent (Python/Groq)	Cross-paper themes, gaps, novel hypothesis generation
Agent 3	Paper Critic (Python/Groq)	Quality score, biases, reproducibility, improvements
Backend	Python 3.9 + FastAPI + Uvicorn	Fast async API; excellent Groq AsyncGroq SDK support
Frontend	HTML5 + Canvas API + Three.js + GSAP	No framework overhead; 3D neural graph; runs in any browser
Fonts	Space Grotesk + Inter + JetBrains Mono	Modern, technical, readable across all devices
Deployment	Docker + Railway.app free tier	Zero-ops for MVP; \$0/month infrastructure

3.5 Risk Register

Risk	Type	Prob.	Impact	Response Plan
<u>Groq API rate limiting at peak usage</u>	<u>Technical</u>	<u>Medium</u>	<u>High</u>	<u>Implement request queue + fallback to local Ollama; cache frequent queries</u>
<u>LLM hallucinations in hypothesis output</u>	<u>Technical</u>	<u>Medium</u>	<u>Medium</u>	<u>Add source-grounding in prompts; display confidence scores; user flagging</u>
<u>Competitor (Elsevier) builds similar tool</u>	<u>Business</u>	<u>Low</u>	<u>High</u>	<u>Accelerate community features; build switching costs via graph data ownership</u>
<u>GDPR violation — storing paper text</u>	<u>Legal</u>	<u>Low</u>	<u>High</u>	<u>Process text client-side; store only metadata; EU server option in v1.2</u>
<u>Key developer leaves the team</u>	<u>Project</u>	<u>Low</u>	<u>Medium</u>	<u>Document all architecture; pair programming; reduce bus factor</u>
<u>Groq changes free-tier pricing</u>	<u>Business</u>	<u>Medium</u>	<u>Medium</u>	<u>Maintain Ollama (local LLM) as zero-cost fallback at all times</u>

4. Planning

4.1 Team Members

Name	Role	Primary Responsibilities	Sprint Focus
Teodor Mihailescu	Lead Dev & AI Architect	FastAPI backend, all 3 agents, frontend panels, D3 graph, Groq integration	WP1, WP2, WP3
Adela Danescu	Backend Dev & QA	API testing, JSON schema validation, bug reports, integration QA, cross-file consistency checks	WP1 (T1.5), WP2 (T2.5), WP3 (T3.7)
Anastasia Sandu	Business & Product	Requirements analysis, SWOT, Lean Canvas, market research, business document authoring	WP1 (T1.4), WP4 (T4.1)
Selena Hurloi	Design & Pitch	UI/UX theme direction, landing page visual design, pitch deck, branding	WP3 (T3.6), WP4 (T4.2, T4.3)

4.2 RACI Matrix

R = Responsible · A = Accountable · C = Consulted · I = Informed

Deliverable / Decision	Teodor M.	Adela D.	Anastasia S.	Selena H.
FastAPI Backend Architecture	R/A	I	I	I
AI Agent Prompt Design	R/A	C	I	I
Frontend & Knowledge Graph	R/A	C	I	C
API Endpoint Testing	C	R/A	I	I
Business Foundation DOCX	C	C	R/A	I
Pitch Deck PPTX	I	I	C	R/A
Landing Page	R	I	I	R/A
Final QA & Submission	A	R	C	C

4.3 Work Packages

The project is structured into four Work Packages (WP), each corresponding to one Sprint.

WP1 — Foundation & Architecture (Sprint 1, Weeks 1–2) | Owner: Teodor Mihailescu

Task	Description	Owner	Status	Acceptance Criteria
T1.1	Repo & dev environment setup	Teodor	Done	All dependencies installed, uvicorn running
T1.2	FastAPI project scaffold	Teodor	Done	main.py with CORS, health endpoint, static serve
T1.3	Groq API integration	Teodor	Done	AsyncGroq client configured, test call succeeds
T1.4	Requirements analysis	Anastasia	Done	MPA 2026 criteria mapped to deliverables
T1.5	Initial README	Adela	Done	Setup instructions documented

WP2 — AI Agent Development (Sprint 2, Weeks 3–4) | Owner: Teodor Mihailescu

Task	Description	Owner	Status	Acceptance Criteria
T2.1	Agent 1 — Paper Analyzer	Teodor	Done	Extracts 9-field JSON; temperature 0.15
T2.2	Agent 2 — Synthesis Agent	Teodor	Done	Returns themes, connections, gaps, 3 hypotheses
T2.3	Agent 3 — Paper Critic	Teodor	Done	Returns 5-dimension score + strengths/weaknesses
T2.4	Chat Agent (bonus)	Teodor	Done	Q&A grounded in full paper library

T2.5	Unit testing of JSON schemas	Adela	Done	All agents return valid JSON under edge cases
------	------------------------------	-------	------	---

WP3 — Frontend & Integration (Sprint 3, Weeks 5–6) | Owner: Teodor / Selena

Task	Description	Owner	Status	Acceptance Criteria
T3.1	D3.js Knowledge Graph	Teodor	Done	Force-directed nodes with concept connections
T3.2	Synthesis panel + render	Teodor	Done	Themes, connections, hypotheses rendered as cards
T3.3	Critique panel + score bars	Teodor	Done	Score breakdown with animated bars, per-paper retention
T3.4	Chat panel	Teodor	Done	Message stream, grounded responses
T3.5	Demo mode (offline)	Teodor	Done	Pre-loaded papers + demo critique data — works without backend
T3.6	UI/UX polish & theme	Selena	Done	Dark purple/cyan theme, topbar agents, toast notifications
T3.7	Integration QA	Adela	Done	All endpoints tested end-to-end in browser

WP4 — Business Docs & Pitch (Sprint 4, Weeks 7–8) | Owner: Anastasia / Selena

Task	Description	Owner	Status	Acceptance Criteria
T4.1	Business Foundation DOCX	Anastasia	Done	SWOT, Gantt, Lean Canvas, Risk Register, WPs
T4.2	Pitch Deck PPTX (11 slides)	Selena	Done	Dark theme, all 4 members, no invented claims
T4.3	Landing page (Three.js+GSAP)	Selena	Done	Accurate features, team section, deliverables links
T4.4	Cross-file uniformisation	Adela	Done	All files consistent: 4 members, 3 agents, no fake data
T4.5	Final submission prep	All	Done	All deliverables linked from landing page

4.4 Sprint Plan

The project follows an Agile methodology with 4 two-week sprints.

Sprint	Details
Sprint 1 Foundation (Weeks 1–2)	Goal: Running FastAPI server with Groq integration, health endpoint live. WP: WP1 DoD: uvicorn starts without errors; /api/health returns 200; all team members can run the project locally.
Sprint 2 Agent Dev (Weeks 3–4)	Goal: All three AI agents implemented, returning valid JSON, tested under normal and edge-case inputs. WP: WP2 DoD: POST /api/analyze, /api/synthesize, /api/critique all return schema-valid JSON; Adela signs off on test cases.
Sprint 3 Integration (Weeks 5–6)	Goal: Full frontend working in browser — graph, synthesis, critique, chat panels integrated and functional in demo mode. WP: WP3 DoD: Demo mode loads 3 papers without backend; all tabs render correctly; critique retained per paper; no JS console errors.
Sprint 4 Delivery (Weeks 7–8)	Goal: All MPA deliverables complete: DOCX, PPTX, landing page accurate and linked. No invented claims. All 4 team members represented consistently. WP: WP4 DoD: Landing page links to all 4 deliverables; business DOCX passes validation; pitch deck reviewed by all team members; submission package ready.

4.5 Gantt Chart

Delivery timeline across 8 weeks. Filled cells (■) indicate active weeks.

Task / Week →	W1	W2	W3	W4	W5	W6	W7	W8
Sprint	S1	S1	S2	S2	S3	S3	S4	S4
Requirements Analysis	■							
FastAPI Scaffold + Groq	■	■						
Agent 1 — Paper Analyzer			■					
Agent 2 — Synthesis				■				
Agent 3 — Paper Critic				■	■			
D3.js Knowledge Graph					■			
Frontend Panels + Chat					■	■		
Integration & Demo Mode						■		
Landing Page							■	
Business DOCX							■	■
Pitch Deck PPTX								■
QA + Submission								■

5. Costs

5.1 Cost Categories & Estimates

Cost Category	MVP (2 months)	Year 1	Notes
Personnel / Salaries	€0	€0	Academic student project — no compensation; team members are unpaid students
Groq API (free tier → Pro at scale)	€0	€1,200	Free up to 14,400 req/day
Cloud hosting (Railway free → VPS)	€0	€240	Free tier for MVP
Domain + SSL	€15	€15	researchmind.ai or similar
Design tools & licences	€0	€240	Figma free tier for MVP
Marketing (social + content)	€0	€600	Organic growth in Year 0
Legal & incorporation	€0	€800	Deferred to post-funding
Total	€15	€3,095	

5.1.1 Token Cost Economics — Groq Inference Pricing

ResearchMind uses Groq Cloud to run LLaMA 3.1 8B Instant. Understanding token costs is essential for financial planning: each agent call consumes a predictable number of tokens, and Groq publishes per-token pricing for paid plans. The free tier (14,400 requests/day) supports MVP scale; the paid tier scales linearly with usage.

Agent / Operation	Input tokens	Output tokens	Cost @ paid tier	Notes
Agent 1 — Paper Analyzer	~2,000	~400	€0.00013	Per paper analyzed (9,000 char input)
Agent 2 — Synthesis Agent	~3,200	~700	€0.00022	Grows with library size; avg 5 papers
Agent 3 — Paper Critic	~1,500	~500	€0.00011	Per paper critique request
Chat Agent	~2,000	~300	€0.00011	Per chat message with library context
Full workflow (1 paper + synth + critique)	~8,700	~1,900	€0.00057	< €0.001 per complete research session

Pricing basis: Groq llama-3.1-8b-instant paid tier = \$0.05/1M input tokens, \$0.08/1M output tokens (as of Q1 2025). Token-to-character ratio: ~4 chars/token. At 1,000 monthly active users performing 20 sessions each: total token cost ≈ €11.40/month. At 10,000 MAU: ≈ €114/month — well within gross margin at Scholar (€19/mo) pricing. Token costs are therefore negligible at current scale and remain below 1% of revenue at Year 1 projections. The free Groq tier (14,400 req/day × 30 = 432,000 req/month) covers approximately 21,600 full workflows, making it viable to run the MVP at zero inference cost.

Scale (MAU sessions/mo) ×	Monthly workflows	Token cost (€)	Groq free tier covers?	Inference % of revenue
100 users × 20 sessions	2,000	€1.14	✓ Yes (free)	<0.1%
1,000 users × 20 sessions	20,000	€11.40	✓ Yes (free)	~0.3%
5,000 users × 20 sessions	100,000	€57	✗ Paid tier needed	~0.6%
20,000 users × 20 sessions	400,000	€228	✗ Paid tier	<1%

5.2 Revenue Model & Projections

ResearchMind uses a freemium SaaS model with four tiers. The free tier drives viral adoption in academic communities, converting 5–8% to paid within 90 days (industry benchmark). University licence deals represent the highest-value channel.

Tier	Price/mo	Year 1 Users	MRR	ARR
Researcher (free)	€0	8,000	€0	€0
Scholar (individual)	€29	400	€11,600	€139,200
Lab (team of 10)	€99	50	€4,950	€59,400
University licence	€499	10	€4,990	€59,880
Total (end of Year 1)		8,460	€21,540	€258,480

5.3 Cost/Benefit Analysis (ROI & Payback Period)

Total 2-year investment: approximately €90,000 (Year 1: €43,595 + Year 2 growth costs: €46,405). Projected Year 2 ARR: approximately €780,000 (3× growth from Year 1 — consistent with SaaS benchmarks for academic tools with strong network effects). Return on Investment (ROI) at end of Year 2: approximately 767%. Payback period: approximately 8 months after product

launch. Break-even point: approximately 620 paid subscribers at average €35/month ARPU.

5.4 Go-to-Market & Marketing Strategy

ResearchMind targets a well-defined, reachable community: academic researchers who are already active online, share papers publicly, and respond to peer recommendations. The GTM strategy avoids paid advertising at launch and instead leverages community-led growth, content marketing, and institutional partnerships — channels that build compounding value over time.

5.4.1 Target Personas

Persona	Description	Pain Point	Entry Channel
PhD Student	Years 1–4, reading 10–20 papers/week for literature review	Overwhelmed by volume; missed connections between papers	Reddit r/PhD, Twitter/X, university lab
Postdoc Researcher	Writing grant applications; keeping up with 2–3 fields	No time for manual synthesis; needs citation-ready hypotheses	Academic Twitter/X, department seminars
Research Lab PI	Managing 5–15 students; reviewing 50+ papers per project	Inconsistent quality assessment across lab members	University IT procurement, direct outreach
Science Journalist	Covering multiple fields; needs rapid paper comprehension	Credibility assessment of pre-prints; identifying novelty	ProductHunt, Twitter/X, LinkedIn

5.4.2 Channel Strategy & Launch Plan

Phase	Channel	Tactic	Budget	KPI
Phase 1 (Month 1–2)	Reddit, Twitter/X, ProductHunt	Launch post on r/PhD, r/MachineLearning, r/academia; ProductHunt Day-1 campaign; Twitter/X demo thread with GIF of 3-agent pipeline	€0	500 signups in 30 days
Phase 1 (Month 1–2)	GitHub	Open-source the backend; README with demo GIFs; star-growth from ML community	€0	100 GitHub stars
Phase 2 (Month 3–6)	SEO / Content	Weekly blog: "How to do a literature review with AI", "Best AI tools for PhD students 2025"; target long-tail keywords	€100/mo	2,000 organic sessions/mo
Phase 2 (Month 3–6)	University pilots	Cold email 20 Romanian & EU university library departments; offer free Lab tier in exchange for feedback + case study	€0	3 pilot agreements
Phase 2 (Month 3–6)	LinkedIn	Founder content: weekly insight on AI + academic research; share synthesis demos; academic network amplification	€0	1,000 LinkedIn followers
Phase 3 (Month 7–12)	Paid social (targeted)	LinkedIn Ads to PhD program coordinators & research group managers; €5 CPC budget; retarget free users	€200/mo	Free→Scholar CVR 5%
Phase 3 (Month 7–12)	Conference presence	Poster / demo at ACL, EMNLP, NeurIPS workshops; NLP & AI communities align with target users	€500 one-time	50 qualified leads/event

Phase	Channel	Tactic	Budget	KPI
Phase 3 (Month 7–12)	Partnership / affiliate	Academic partnerships; blogger Zotero cross-promotion; community ResearchRabbit referral exchange	€150/mo	10% of new signups

5.4.3 Marketing Budget & KPIs

Metric / Budget Line	Month 1–2	Month 3–6	Month 7–12	Year 1 Total
Marketing budget (€/mo)	€0	€100	€350	€2,250
Target MAU (monthly active users)	200	1,200	5,000	—
Target registered users	500	3,000	12,000	12,000
Free → Paid conversion target	—	3%	5%	5–8%
Target paying subscribers	—	90	600	600+
NPS target (researcher satisfaction)	—	>40	>55	>55
CAC target (cost per paid user)	€0	<€3	<€5	<€5

6. One-Page Summary — Lean Canvas

PROBLEM Researchers waste 6h/week searching, reading, and connecting papers manually. No tool builds a living AI knowledge graph with automated quality critique.	SOLUTION Three AI agents: (1) analyse any paper in 10 sec; (2) synthesise library and generate hypotheses; (3) critique paper quality and flag weaknesses.	UNIQUE PROPOSITION The only tool that turns your paper library into an AI knowledge graph with multi-agent synthesis, hypothesis generation, AND quality critique.
CUSTOMER SEGMENTS Primary: PhD students & postdocs. Secondary: Academic labs & institutes. Tertiary: R&D in pharma & tech.	CHANNELS Academia Twitter/X · ArXiv community · University partnerships · ProductHunt launch · Research lab pilots	UNFAIR ADVANTAGE Groq-powered inference (700 tok/s, free); open-source LLaMA 3.1; 3-agent architecture no competitor has; academic team with research domain expertise.
COST STRUCTURE MVP (2 months): €3,390 (human capital €3,375 + domain €15) Year 1: ~€43,595 Break-even: Month 8 at 620+ paid subscribers		REVENUE STREAMS Freemium SaaS: Free (€0) Scholar (€29/mo) Lab (€99/mo) University licence (€499/mo) Target ARR Year 1: €258,480

7. AI Usage Disclosure

In accordance with MPA 2026 guidelines, the following AI tools were used in the creation of this project:

Tool	Usage	Example Prompt / Command
Claude (Anthropic)	Full project generation: landing page, MVP code, business document, pitch deck	"Read the MPA 2026 requirements and build a full startup project with latest libraries, animations, and free AI. Rename to ResearchMind."
Groq + LLaMA 3.1	MVP runtime AI: Agent 1 (paper analysis), Agent 2 (synthesis), Agent 3 (critic)	"You are the Paper Analyzer Agent. Extract structured knowledge from this paper as JSON..."